

Project: Increasing Throughput with SHA

Assigned: Monday 10/28; Due **Friday 12/6** (midnight)

Instructor: James E. Stine, Jr.

1. Introduction

This project is meant to be an encompassing project that gives you the full experience of all that you learned in this course. It will bring ideas that we covered as well as had within laboratories. It will also try to reinforce all the skills you learned during your time in laboratory.

This project is significantly involved compared to our normal labs and much of the detail is left up to you to implement. Unfortunately, we cannot monitor your process, therefore, it is important to allocate your time wisely and come into laboratory every week. Do not assume you can do much of the work later or will have time to make up time later. To make things worse, we only have one week when we return from Thanksgiving making it almost impossible to complete the project during that final week. That is, you should complete most of the project items before Fall Break and the Thanksgiving Holiday. Waiting until the final week will be a guaranteed disaster for you.

When designing digital systems, the process of creating logic to accomplish some form of computation is a function of the blocks that can be utilized. Many times, the only way to increase the amount of computation is to process more and more data in parallel. However, sometimes computation needs to be processed independently per block and the only way to increase the computation time is to increase the amount that can happen per time. This is sometimes called throughput.

2. Background

Throughput can be visualized with a kitchen. Suppose that someone only has 1 oven in their kitchen. Obviously, the cook is limited by the amount of time that can complete a dinner by how often I can use this oven. On the other hand, if I am fortunate to have 2 ovens in my kitchen, I can use this to help do more things at the same time. This processing power is predicated on the fact that each oven probably is not used at the same time and some foods cook faster than others. Therefore, sequencing the oven to use it optimally is the key to making throughput work effectively.

In digital systems, one way we can increase the throughput for a design is to add a register where one stores intermediate results. These intermediate results are typically stored into a group of flip flops or a register and fed back into the register to get to the final computation. The use of registers helps optimize this throughput and also increases the latency of the computation by increasing the clock rate.

For example, suppose a computation needs to pass through 4 adders to form the final result and stored into a register. Also, assume that each adder completes its operation in 100 ns and stores the result into a register. In theory, this would be $4 \times 100 = 400$ ns to complete the operation. If a designer removed 3 of these adders and added a register following an adder (note: the result would need to go back into the adder 3 more times to complete the operation as shown in Figure 1). For both of these scenarios, it would still take a little time to store/read the value from a register, but for this example, I am assuming that this amount is negligible. However, the first unit can only be clocked at $1/400\text{ns} = 2.5$ MHz, however, the second version could be computed at $1/100\text{ns} = 10.0$ MHz and run significantly faster. If a company is designing this hardware to process 100,000 operations, the first version can only process all of these items in 40 msec, whereas, the clocked version takes 10 msec. This is a 4x speedup and significant improvement in latency for the total operation. Most importantly, the second architecture may have less area than the first architecture. Typically, more registers can increase energy consumption, however, if you can significantly gain performance by increasing throughput with a small increase in energy, it is often worth the investment. This is what we are investigating for this project or increasing the throughput of our design.

The key to digital systems is that digital logic can process data in parallel. This project involves taking Laboratory 2 and modifying it to include a register to increase throughput and performance. In order to complete this project, the following items must be functional on the board:




Figure 1: Adding registers to increase throughput

- Control logic should be added to control each iteration of the unit to compute the SHA-256 operation.
- All testbenches for the system, control, and datapath.
- Use of the ILA to test your design (see Section 3.1)
- Cleaned up code removing any unnecessary logic and/or states.
- A report documenting the complete system and its operation indicating that it is working.
- Anything you can think of to promote your project!

This project is not difficult and it is an excellent choice for groups that want a straightforward project and/or who wish to spend a minimal amount of time as possible on the project. However, it is in your best interest to make the project better! Any modification beyond the baseline project will incur extra points that could help compensate for a bad test, missing homework, or bad quiz. The following are potential modifications you can easily add for extra credit:

- Add more operations (e.g., SHA-512)
- Add some way to add a new input message to the unit as opposed to hard coding the messages.
- Determine a way to optimize the operation of the register to process throughput faster (e.g., adding another `main_comp` block so that it completes faster and balances the timing between clocking)
- Determine the best timing to optimize the throughput
- Add a poster to explain to the masses!
- Create a YouTube[®] channel on this project.

3. Implementation

The basic idea for the inputs and output includes your move which you should add through the switches. You will technically only need a “Start” signal to get things going in the FSM. As explained in class, you should create the FSM intelligently, so that it can theoretically process more data when needed.

The main key elements of the design will be provided to you, such as the `sha256.sv` module. But, you will have to add registers, control logic (e.g., Finite State Machines), multiplexors, input/output signals. It is advisable to use as many input/output signals to help you debug what is going on as you simulate and get to your final implementation.

You should use many of the elements discussed in our textbook [1]. I would highly advise using the *idioms* we discussed in class and in the textbook for things like registers making sure that you also reset or enable them appropriately. That is, you should also incorporate a **reset** somewhere into your design. All of the HDL discussed in the textbook [1] is on Canvas as a zip file. Some of the registers are also available in the **SV** subdirectory of your project repository. Feel free to use this HDL in any way you wish.

Again, using sequential logic to increase the throughput of a system involves adding stages of registers or pipelines to split complex tasks into smaller, manageable steps that can be processed concurrently. This approach improves throughput by allowing different stages of a task to be processed simultaneously, leading to an overall increase in the number of tasks completed per unit time.

As explained previously there are many advantages in this process, but there are also disadvantages. As explained in class, the key to any engineering system is making sure that you have more advantages than disadvantages for it to be successful. Some advantages to this approach are: increased throughput, reduced latency per task, enhanced resource utilization, and scalability. The disadvantage to this process is that its more challenging and complex than a traditional serial approach. There are other disadvantages, such as minimal increase in latency, but for the most part using this approach saves a ton of time. That is, carefully balancing these trade-offs is key to effectively using sequential logic to optimize throughput.

To make sure your implementation works, you should use some kind of an enabled flip flop (i.e., `flop_enr.sv`) that stores data when needed. The FSM will be utilized to make sure these registers are written at the right time. You can utilize the methodology in class to design the datapath and control of this design, however, any approach that is reasonable also works.

3.1 Integrated Logic Analyzer (ILA)

In Lab 2, we used the 7-segment display to display the final result. Although this works, many designers use Field Programmable Gate Array (FPGA) features for debugging called the Integrated Logic Analyzer (ILA). In the 1970s and 1980s, many logic designers utilized separate units that enabled them to debug items going on in their logic. These Logic Analyzers, such as the Tektronix LV-500, enabled designs to be tested somewhat automatically. However, with the advent of the FPGA, these ideas could be automatically integrated into the FPGA logic. More importantly, with this idea and a little ingenuity, you can theoretically monitor internal signals.

For this laboratory, we are going to use the ILA from AMD. The customizable Integrated Logic Analyzer (ILA) IP core is a logic analyzer core that can be used to monitor the signals of a design. The ILA core includes many advanced features of modern logic analyzers, including Boolean trigger equations, and edge transition triggers. Because the ILA core is synchronous to the design being monitored, all design clock constraints that are applied to your design are also applied to the components inside the ILA core.

These ILAs are not perfect in that they consume onboard memory blocks that the FPGAs have. So, in theory, you cannot probe all the signals in your design. However, not all signals usually need to be probed using an ILA. There is more information in the UG936 AMD (v2024.1) manual in the docs directory. We tried to put together a document in the `debug_tutorial` subdirectory of your repository.

You should decide which signals you would like to see if your design is working. Again, you have the ability to probe a fair number of signals in your FPGA, but be careful not to probe all your signals. You can use the clock of the FPGA as the main driver for this ILA block as documented in our debug document.

4. Tasks

Most of the main modules have been given to you to help you understand the problem better. In fact, we have given much of the project for you in this text. The difference here is that we have not given you any testbenches and you will have to figure out how things go together. You have all the skills you need to complete this project, so trusting yourself and the process is worthwhile and I know all of you can do this. The tasks of the project are as follows:

1. Add registers and any other logic to pipeline your system accordingly from `sha256.sv`. It is easiest to use 1 version of the `main_comp` block and put registers after this unit. You are welcome to determine

the most optimal number of `main_comp` blocks to use, but to get the baseline project done, I would use only 1 block.

2. Complete block diagrams of your system, include detailed interface specifications listing all signals and describing their timing.
3. Design a control logic presumably with a FSM to have it work correctly.
4. Build the testbench to simulate your design and make sure things are working for both the datapath, control and combined datapath/control. It is advisable to make the datapath work first by creating a testbench that makes your datapath run. Then, add the control logic to make sure your datapath works as needed.
5. Once your design completely works, implement the design on the DSDB board. You can still use the switches, push buttons, LEDs and also use the seven segment display to output key elements of your project. Remember, to use the LEDs to help you debug your design.
6. Document the output of ILA for the message and hashed output. You can use the ILA to debug your design, as well. This should be significantly easier than using other items (e.g., 7-segment display). Once you have the ILA set up, it is conceivably easy to run over and over.
7. Finally compare the PPA of both designs and provide a justification if this project's design can process 1 million SHA vectors faster or slower than Lab 2's implementation. Please be aware that the area of your ILA-enabled design may be larger due to the added items needed to support the ILA. The resource output should actually document the ILA separately.

Again, the process here is not difficult. If you need to work out any of the procedures or ask me to inspect your design, I would recommend stopping by to ask questions or advice. I would not advise waiting until the last week to start as I might be busy with end-of-the-semester duties and starting early is the best practice.

4.1 Extra Credit

There are lots of opportunities for extra credit with this project. But, please, first focus on completing the baseline project before attempting the extra credit option. One of the advantages of digital logic is that many bits can be computed in parallel and then chosen later to be correct or incorrect.

Because Field Programmable Gate Arrays (FPGAs) can contain many millions of logic gates, you can do lots of optimization for this project. Timing is challenging in this lab and making sure your design works for multiple messages requires some thought in how to process things accordingly.

5. Video and Lab Report

I am asking for both a final report and video demo of your design. You can easily create a video on your cell phone that is no more than 5-10 minutes that encapsulates your design and how it works. Please work consistently throughout the final weeks of the semester to make sure you complete the project on time.

I will also give extra credit to those that put a little effort into making an outstanding video and showcasing their project in detail. You could also potentially discuss other topics including significant additions to your project.

You are also required to submit a final report of your design using the lab rubric. You should remember to submit both your lab report and video report to Canvas for your team, but please also submit your team evaluation, as well. Beware; no late projects will be accepted and if you miss submitting your project on time, you will receive a 0 for your project grade! This procedure should be similar to what you are using for your labs. You should also take a printout of your waveform from your ModelSim simulation. Only one of your team members should upload the files, lab report, and team assessment. Also, please make sure you hand in all files, including your HDL, testbenches, and other important files you wish for us to see.

For the final project, you cannot use “oops” days to help you. The project is due December 6, 2024 at midnight. If you are running late, please hand in something. If you decide not to hand in anything, you will receive a 0 for the project. Only medical or family emergencies that are clearly documented for getting any extensions. Even if you get an extension, it will be difficult as the University closes the Endeavor building early. Also, please be aware that Thanksgiving week falls at a bad time this year. Typically, you have two weeks after Thanksgiving before the semester ends. Unfortunately, this year only has one additional week and you will probably not get anything done then. I highly recommend trying to get as much done before the Thanksgiving break to avoid any issues.

Please contact James Stine (james.stine@okstate.edu) for more help. Your code should be readable and well-documented. In addition, please turn in additional test cases or any other added item that you used. Please also remember to document everything in your Lab Report using the information found in the Grading Rubric.

References

- [1] Sarah Harris and David Harris, *Digital Design and Computer Architecture: RISC-V Edition*, Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1st edition, 2021.